IEEE 37th International Conference on Distributed Computing Systems
Atlanta, GA, USA

MELODY: A Long-time Dynamic Quality-aware Incentive Mechanism for Crowdsourcing

Hongwei Wang[1, 2], Song Guo[2], Jiannong Cao[2], Minyi Guo[1]

[1] Dept. of Computer Science and Engineering, Shanghai Jiao Tong University

[2] Dept. of Computing, The Hong Kong Polytechnic University

06/08/2017

上海交通大學
SHANGHAI JIAO TONG UNIVERSITY

THE HONG KONG
POLYTECHNIC UNIVERSITY
香港理工大學

Background (1/3)

- **Crowdsourcing** allows requesters to allocate tasks to a group of workers on the Internet to make user of their collective intelligence

- Crowdsourcing tasks are typically …
  - small
  - quite difficult or too expensive to automate
  - simple for humans
  - e.g., proofreading and image labeling

Background (2/3)

- AMT (Amazon Mechanical Turk)

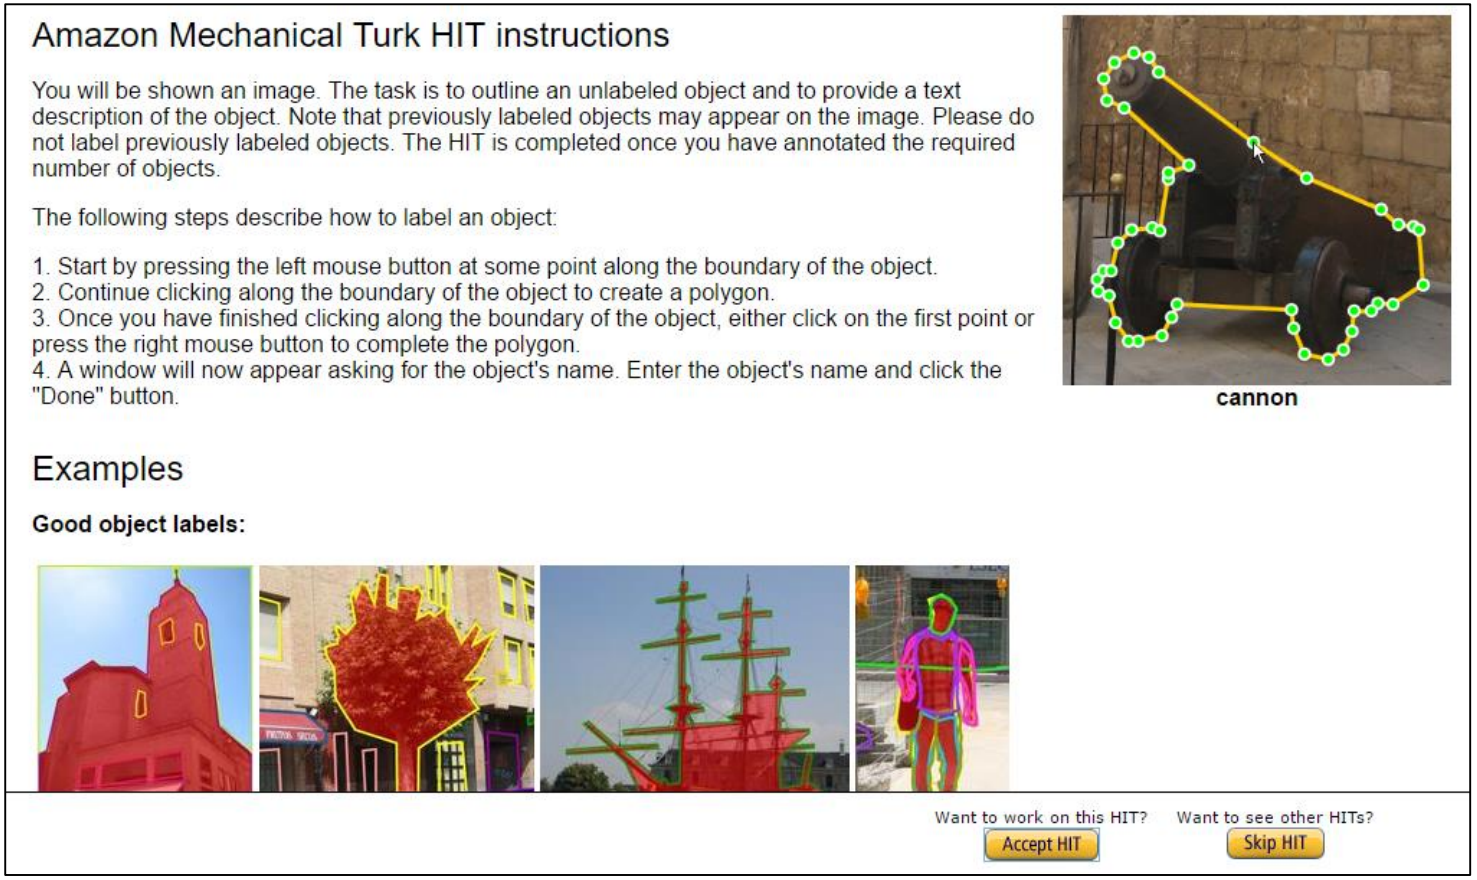

Background (3/3)

- **Quality Control** is one of the key considerations in crowdsourcing
- Quality Control is nontrivial because …
  - crowd workers are at different levels of problem-solving capability, and may vary over time
  - it's hard to incorporate quality control naturally into designs of crowdsourcing platforms

We try to solve the above two challenging issues in this work

Object to be outlined

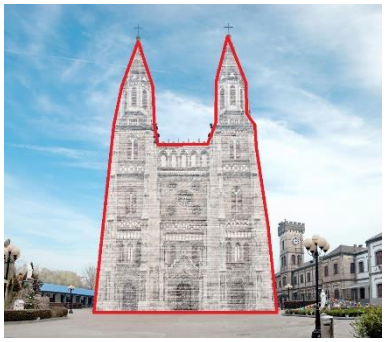

Good job

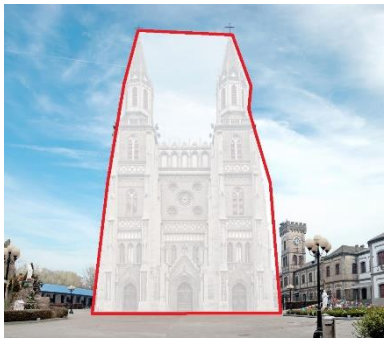

Bad job

Motivation (1/3)

- Existing incentive mechanisms for crowdsourcing are …

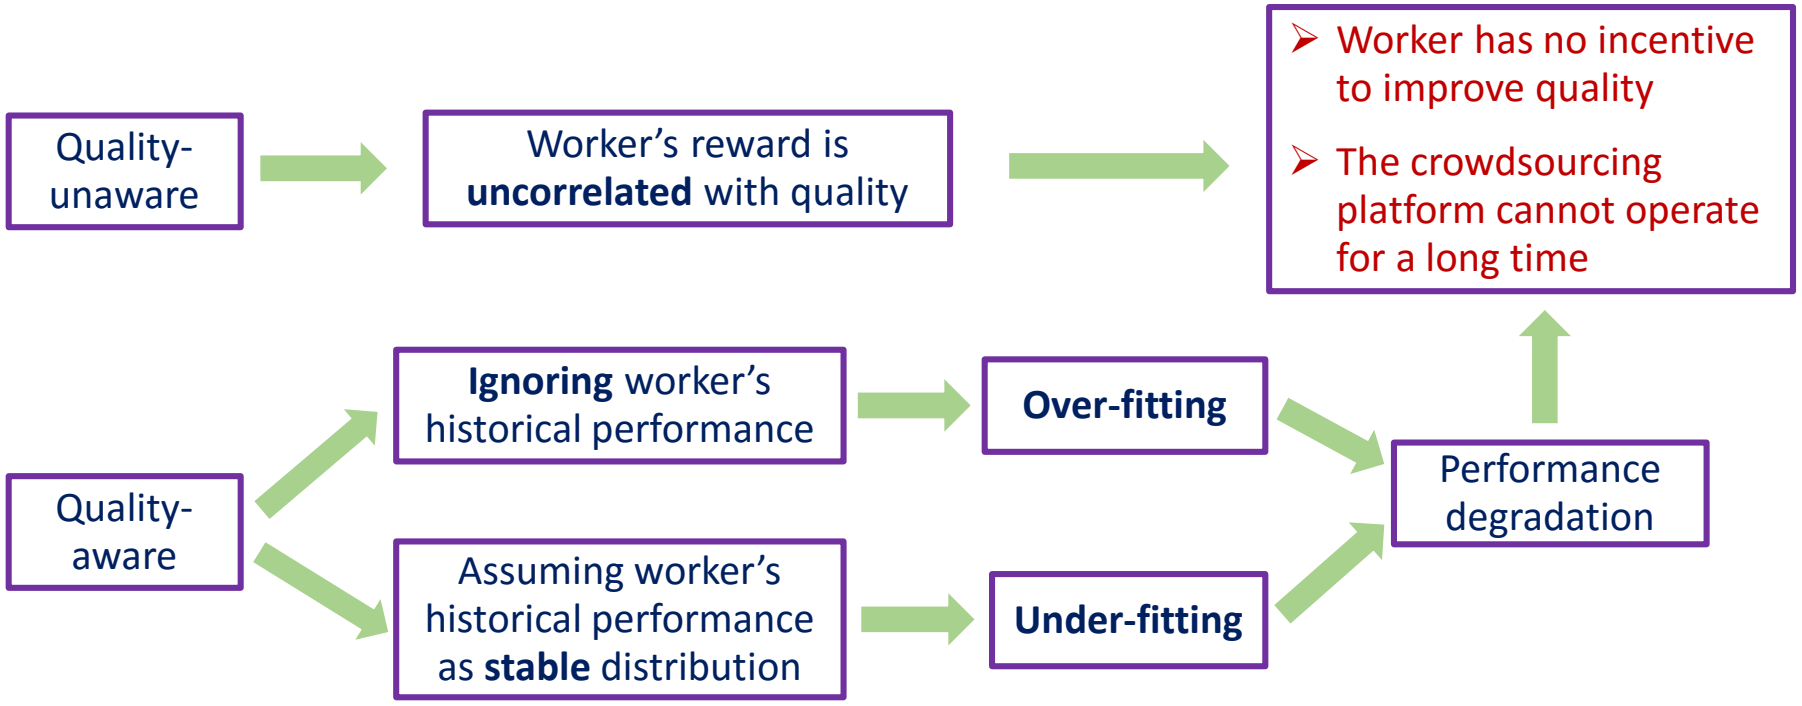

Motivation (2/3)

What is the best prediction of worker's quality at time t based on his observed historical performance?

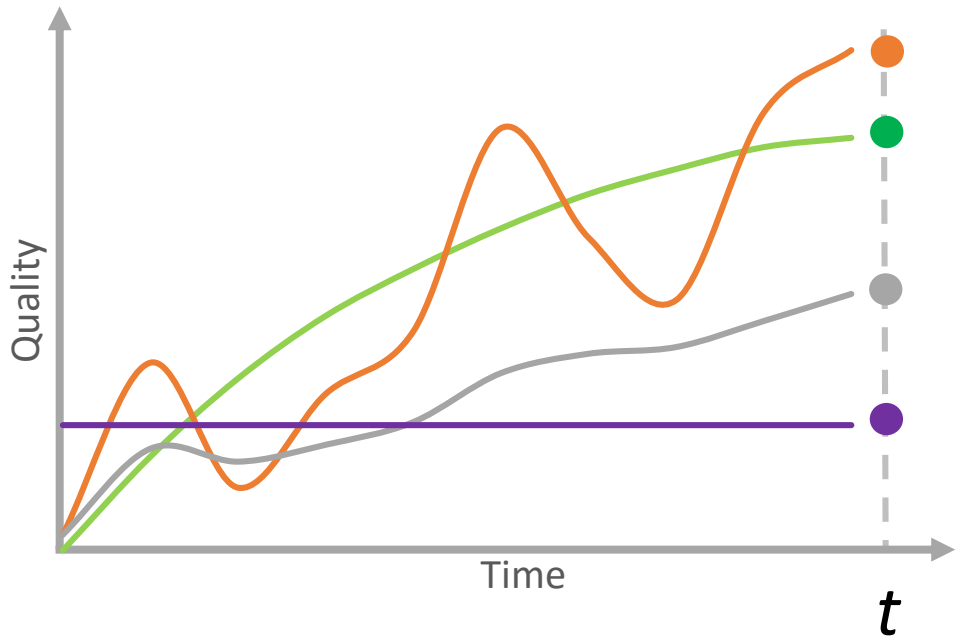

Observed performance
(over-fitting prediction)

Latent quality
(optimal prediction)

Cumulative average
(under-fitting prediction)

Fixed value
(Quality-unaware prediction)

Motivation (3/3)

Can we really predict workers' quality based on their observed historical performance?

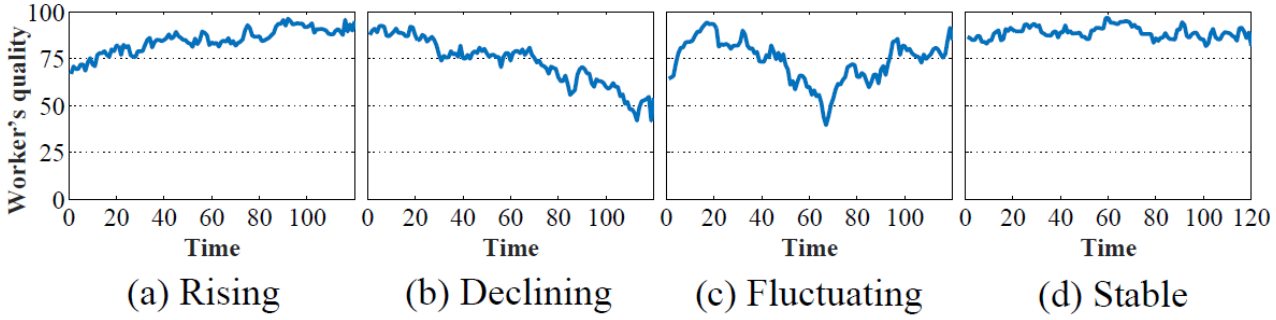

Four typical types of workers' long-term quality curves in reality. The data come from crowdsourcing tasks conducted on AMT for affective text analysis. We measure workers' quality as "maximum rating - average error", where the average error is the difference between ground truth and worker's average rating over 10 items.

MELODY

So what is MELODY?

MELODY is an incentive **ME**chanism considering workers' **LO**ng-term **DY**namic quality for crowdsourcing markets which satisfies properties of:

- Truthfulness
- Individual rationality
- Competitiveness
- Computational efficiency
- Budget feasibility
- Long-term quality awareness

System Workflow in One Run

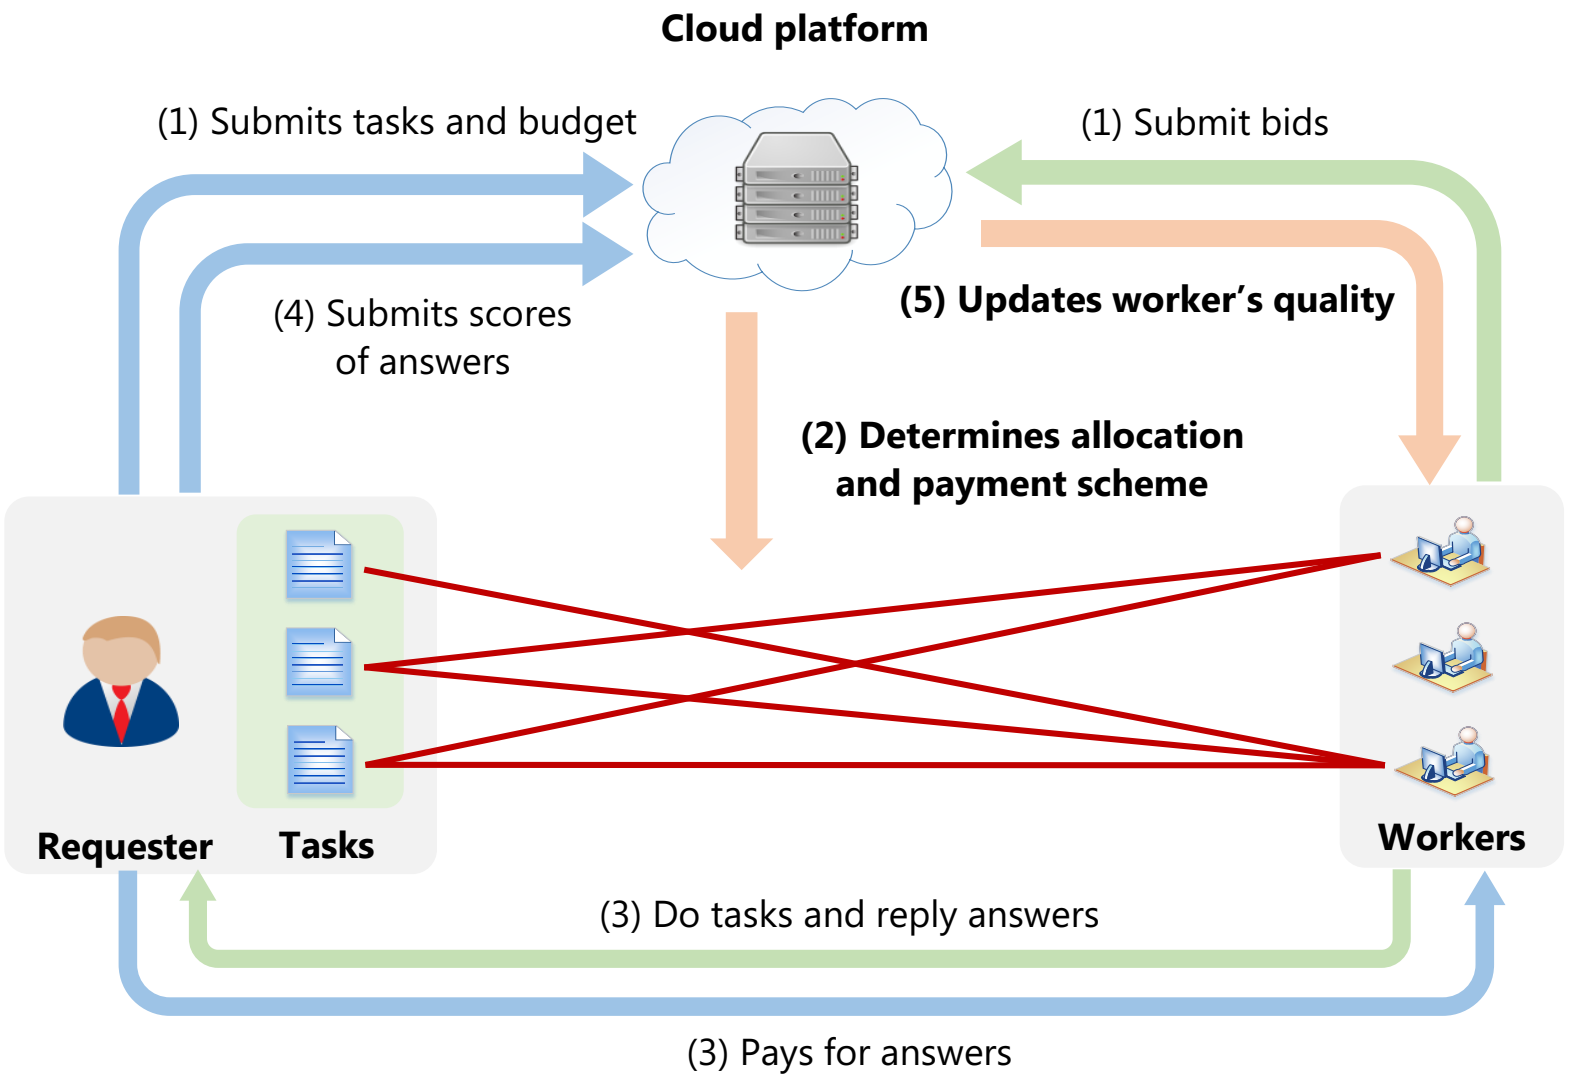

Worker Modeling

- During task allocation in run r, each worker i associates:
  - A bid of **cost** for performing every single task: c_i^r (may not be true)

It costs me **2** yuan for each task in this run …

Worker Modeling

- During task allocation in run r, each worker i associates:
  - A bid of **cost** for performing every single task: $\boldsymbol{c_i^r}$ (may not be true)
  - A bid of maximum number of tasks (also call **frequency**) he is willing to complete: $\boldsymbol{n_i^r}$ (may not be true)

It costs me **2** yuan for each task in this run …

I'd like to do **4** tasks at most in this run…

Worker Modeling

- During task allocation in run r, each worker i associates:
  - A bid of **cost** for performing every single task: $\boldsymbol{c_i^r}$ (may not be true)
  - A bid of maximum number of tasks (also call **frequency**) he is willing to complete: $\boldsymbol{n_i^r}$ (may not be true)
  - A **quality index** given by the platform: $\boldsymbol{\mu_i^r}$

It costs me **2** yuan for each task in this run …

I'd like to do **4** tasks at most in this run…

Quality index: **4.2** / 5

Utility of Worker and Requester (1/3)

- **A worker *i*'s utility** is the difference between total payment he receives and his total cost

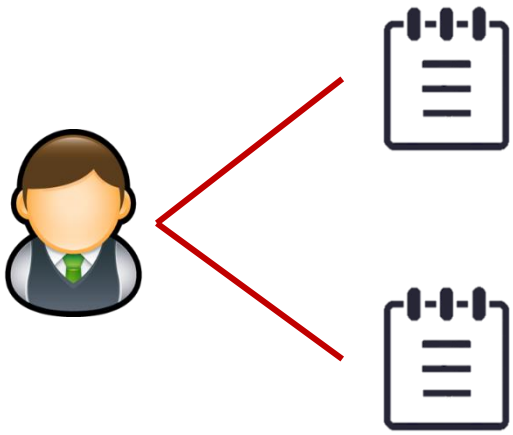

payment		cost		utility per task
4	-	2	=	2
3	-	2	=	1

Utility of this user: 2 + 1 = 3

Utility of Worker and Requester (2/3)

- A task j is **satisfied** if the integrated quality received by this task j is greater than a threshold Q

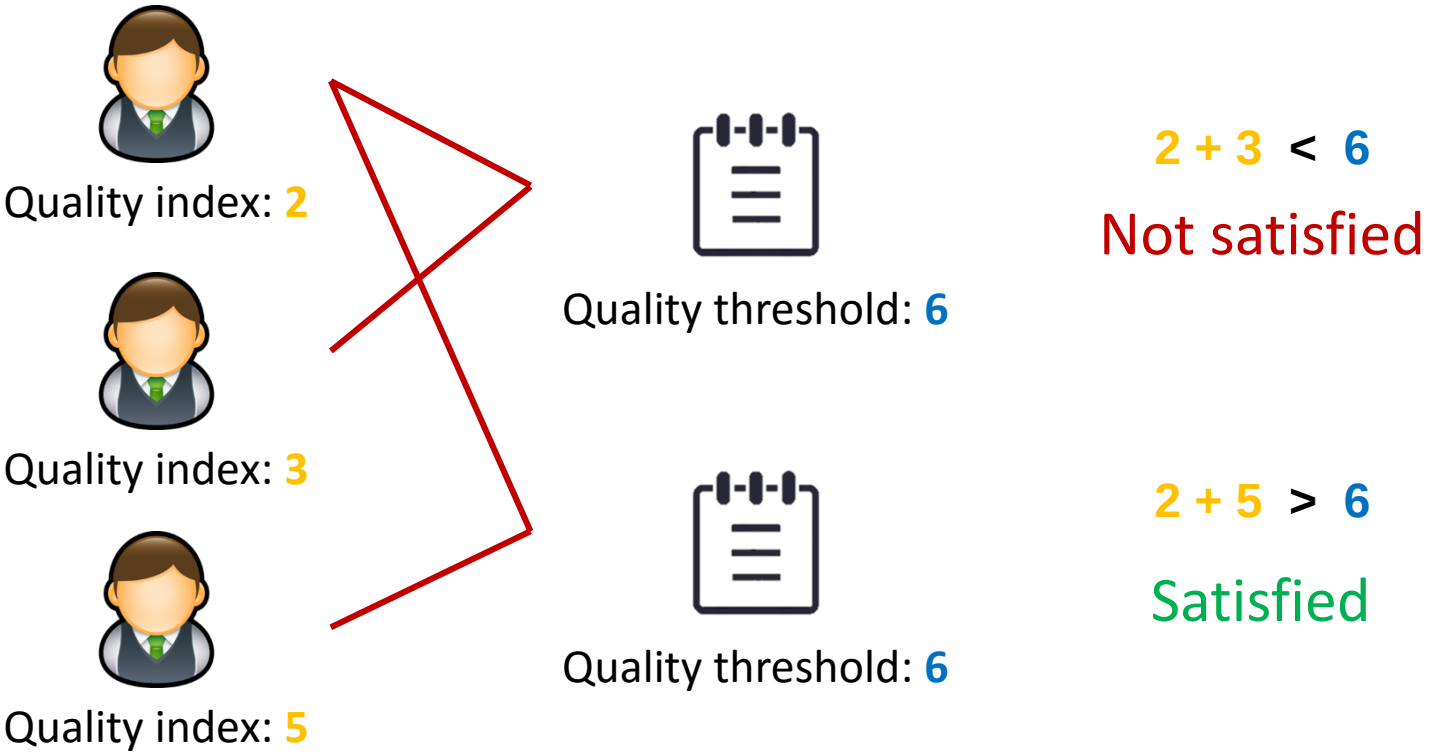

Utility of Worker and Requester (3/3)

- **A requester's utility** is the number of satisfied tasks

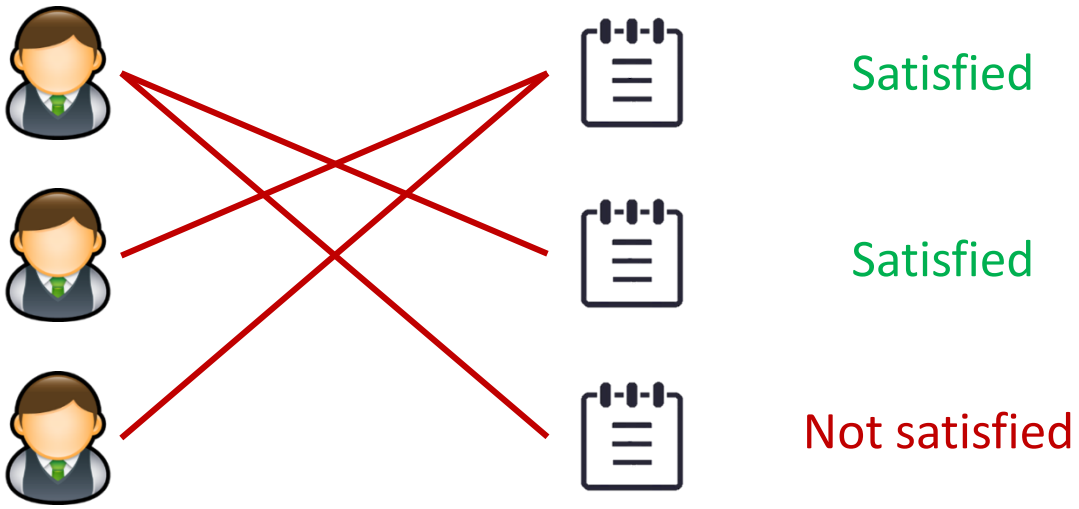

Utility of the requester: 2

Design Objectives

- The **Single Run Auction (SRA)** problem (NP-hard):

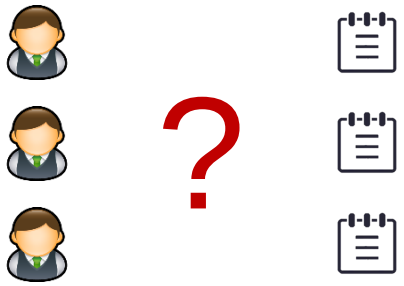

$$
\begin{aligned}
\max \quad & U^r \\
\text{s.t.} \quad & \sum_{i \in \mathcal{W}^r, t_j^r \in \mathcal{T}^r} p_{i,j}^r \leq B^r \\
& \sum_{t_j^r \in \mathcal{T}^r} x_{i,j}^r \leq n_i^r, && \forall i \in \mathcal{W}^r \\
& x_{i,j}^r \in \{0,1\}, && \forall i \in \mathcal{W}^r, \forall t_j^r \in \mathcal{T}^r
\end{aligned}
$$

- We hope to design MELODY satisfying …
  - **Truthfulness** (workers won't lie about their bids)
  - **Individual rationality** (workers' utilities are always non-negative)
  - **Competitiveness** (MELODY's performance is close to optimal solution)
  - **Computational efficiency** (MELODY runs within polynomial time)
  - **Budget feasibility** (Budget constraint is hold)
  - **Long-term quality awareness** (MELODY can predict workers' future quality according to its long-term characteristics)

MELODY Design for SRA (1/2)

- **Step 1**: sort all workers in descending order of μ_i/c_i;
- **Step 2**: sort all tasks in ascending order of Q_j;
- **Step 3**: Allocate available worker i to task t_j in the corresponding order, and pay $\frac{c_{k+1}}{\mu_{k+1}}\mu_i$ for the allocation; ($k+1$ is the first worker who does not get task t_j)
- **Step 4**: Calculate the total payment for completing each task, and select as many tasks as possible under the given budget.

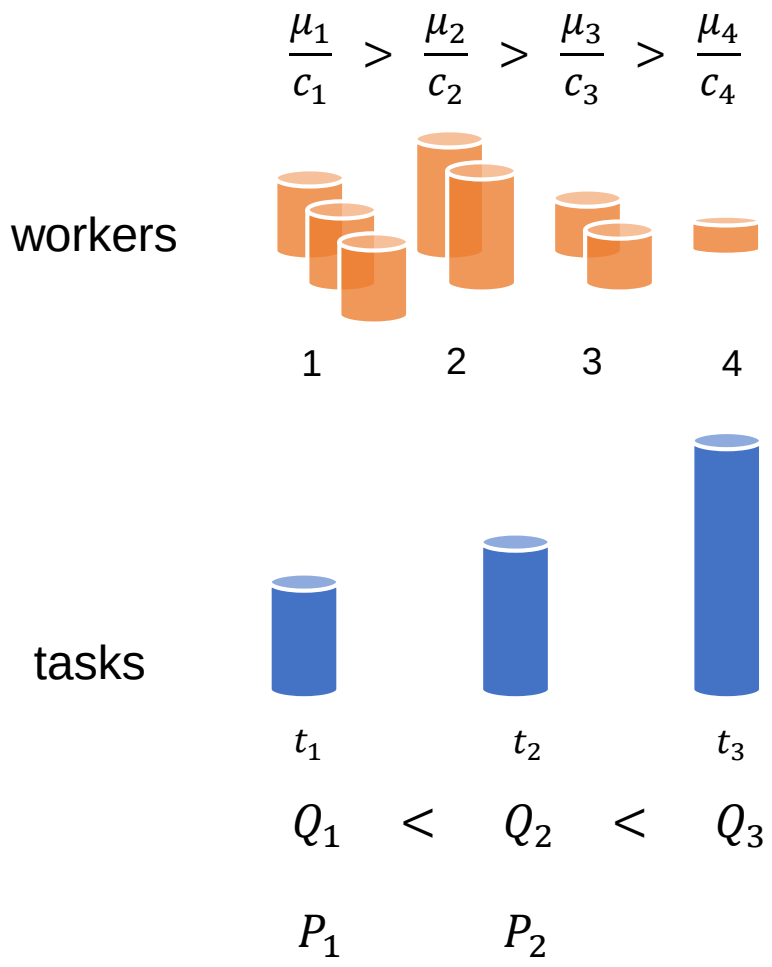

MELODY Design for SRA (2/2)

- **Step 1**: sort all workers in descending order of μ_i/c_i;
- **Step 2**: sort all tasks in ascending order of Q_j;
- **Step 3**: Allocate available worker i to task t_j in the corresponding order, and pay $\frac{c_{k+1}}{\mu_{k+1}}\mu_i$ for the allocation; ($k+1$ is the first worker who does not get task t_j)
- **Step 4**: Calculate the total payment for completing each task, and select as many tasks as possible under the given budget.

- ✓ **Truthfulness**
- ✓ **Individual rationality**
- ✓ **Competitiveness**
- ✓ **Computational efficiency**
- ✓ **Budget feasibility**

Algorithm 1 MELODY Design for the SRA Problem

Input: $\mathcal{W}_U$, $\mathcal{T}$, B;
Output: $\mathcal{X}$, $\mathcal{P}$;
1: $\mathcal{W} \leftarrow \{i \in \mathcal{W}_U | \Theta_m \leq \mu_i \leq \Theta_M \text{ and } C_m \leq c_i \leq C_M\}$
2: Sort all $i \in \mathcal{W}$ in descending order of μ_i/c_i;
3: Sort all $t_j \in \mathcal{T}$ in ascending order of Q_j;
4: $x_{i,j} \leftarrow 0$, $p_{i,j} \leftarrow 0$, $P_j \leftarrow 0$ **for each** $i \in \mathcal{W}$, $t_j \in \mathcal{T}$;
5: **for all** $t_j \in \mathcal{T}$ **do**
6: Find the smallest k such that $\sum_{i \leq k \text{ and } n_i > 0} \mu_i \geq Q_j$;
7: **if** such k exists **then**
8: **for all** i s.t. $i \leq k$ *and* $n_i > 0$ **do**
9: $x_{i,j} \leftarrow 1$;
10: $p_{i,j} \leftarrow \frac{c_{k+1}}{\mu_{k+1}}\mu_i$;
11: $n_i \leftarrow n_i - 1$, $P_j \leftarrow P_j + p_{i,j}$;
12: **end for**
13: **end if**
14: **end for**
15: $\mathcal{X} \leftarrow \emptyset$, $\mathcal{P} \leftarrow \emptyset$;
16: Remove all zero P_j and sort the remaining in ascending order;
17: **for all** P_j s.t. $P_j \leq B$ **do**
18: $\mathcal{X} \leftarrow \mathcal{X} \cup \{x_{i,j} \mid x_{i,j} = 1, i \in \mathcal{W}\}$;
19: $\mathcal{P} \leftarrow \mathcal{P} \cup \{p_{i,j} \mid p_{i,j} > 0, i \in \mathcal{W}\}$;
20: $B \leftarrow B - P_j$;
21: **end for**
22: **return** $\mathcal{X}$, $\mathcal{P}$;

System Workflow in One Run

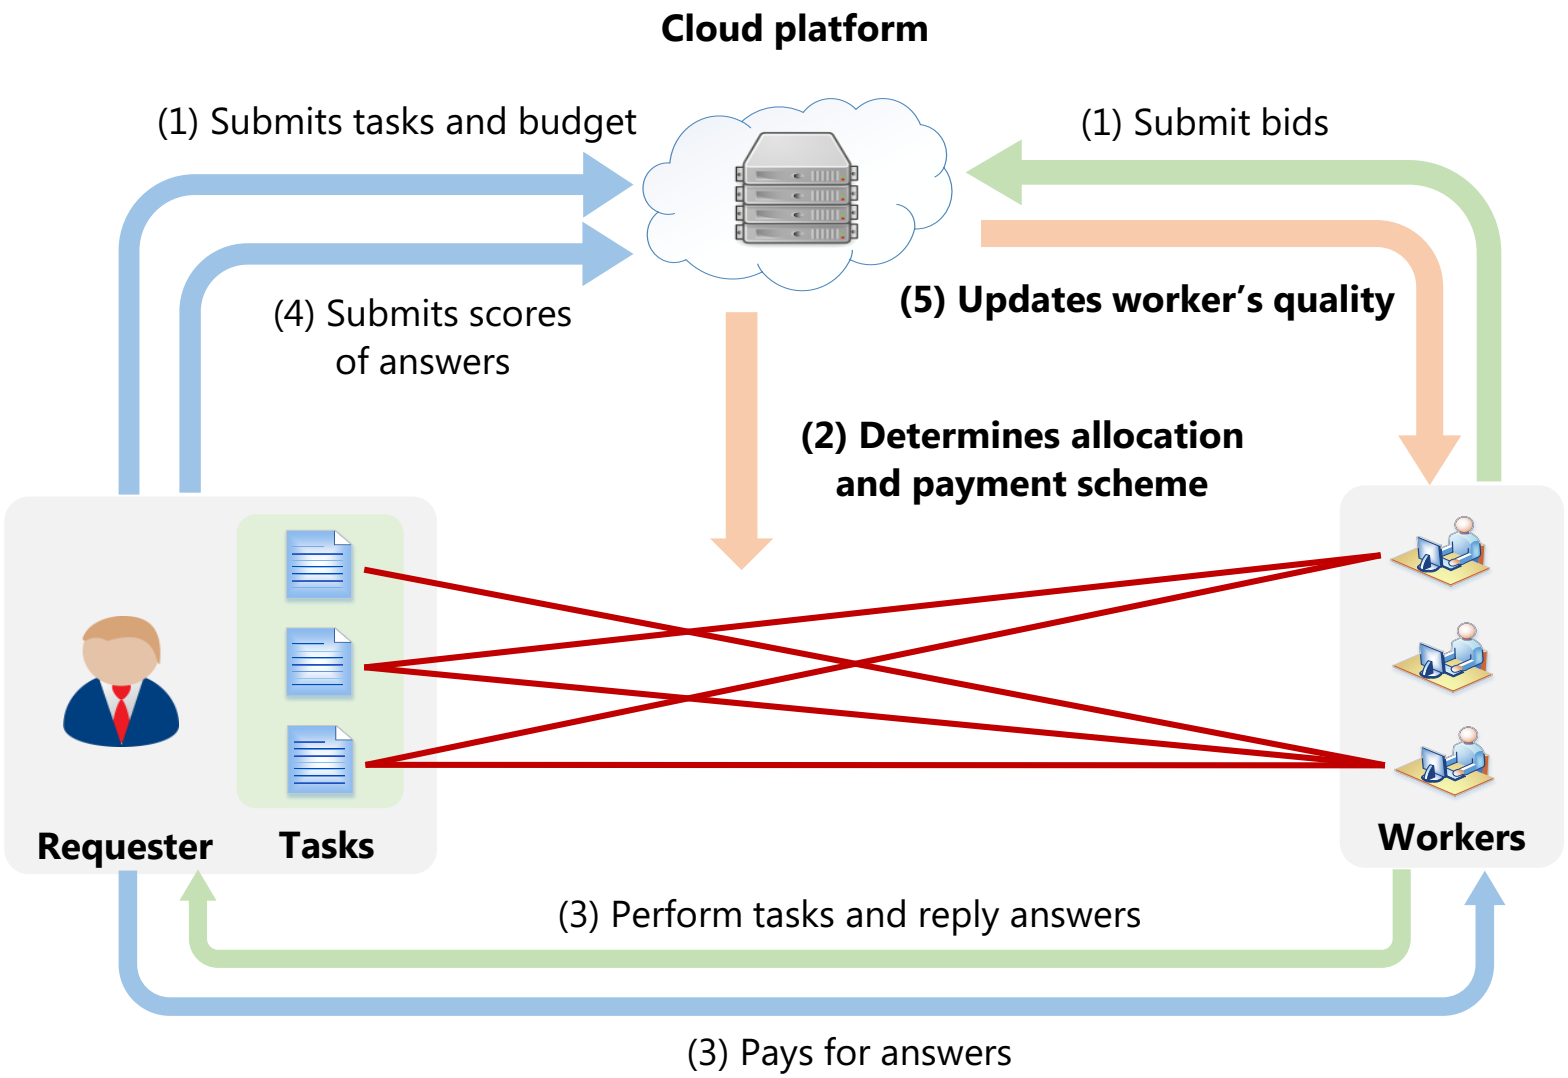

MELODY Design for Quality Updating (1/3)

Over-fitting and under-fitting

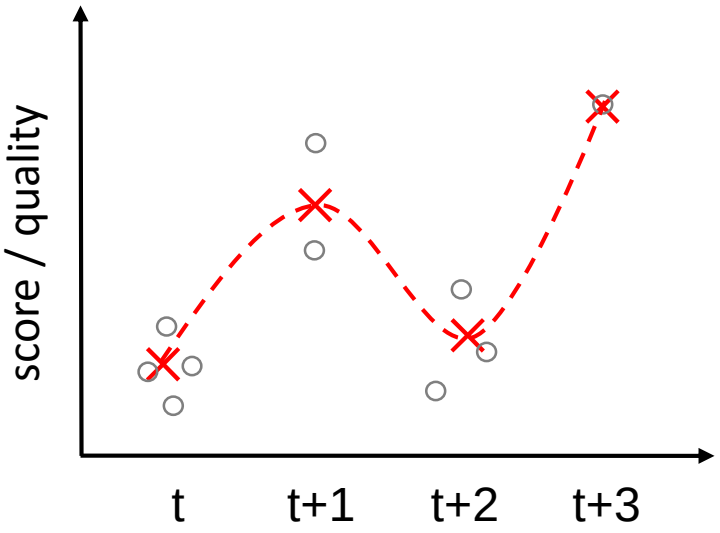

Per-run average

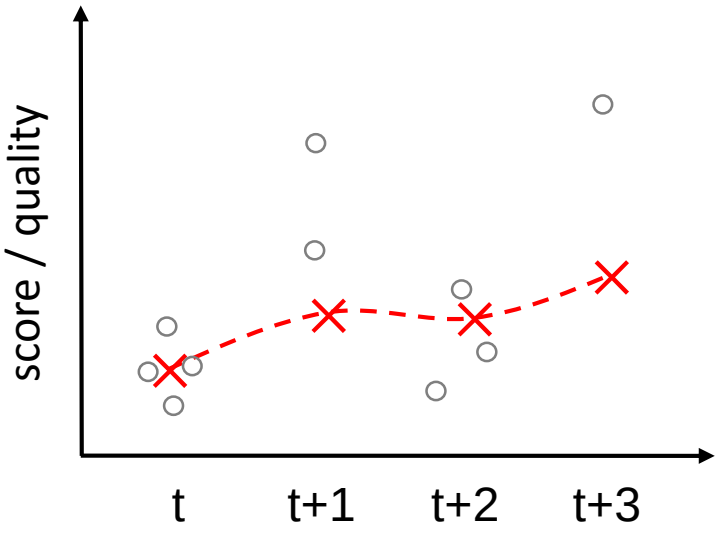

Cumulative average

○ Observed scores

× Estimated quality

MELODY Design for Quality Updating (2/3)

Linear dynamical systems

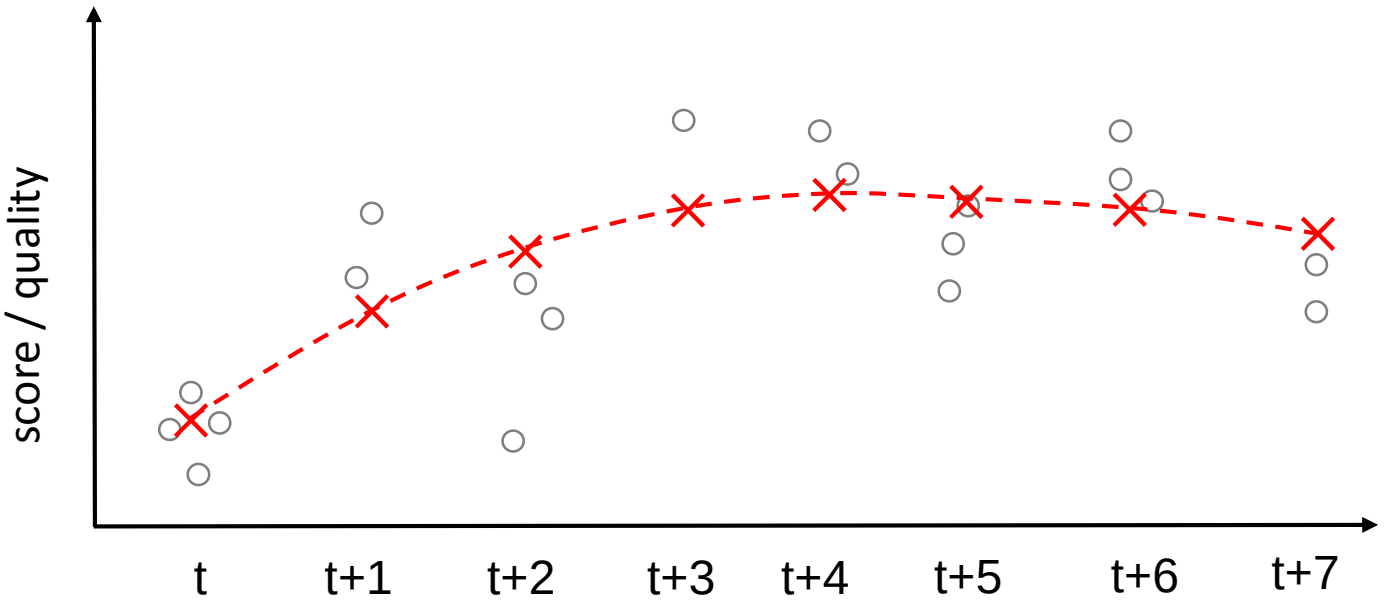

MELODY Design for Quality Updating (3/3)

- Hyper-parameter learning:
  Expectation Maximization (EM)
  - Complete-data likelihood function:

$$L(\mathbf{S^r}, \mathbf{Q^r}; \boldsymbol{\theta}) = \log p(\mathbf{S^r}, \mathbf{Q^r}; \boldsymbol{\theta})$$
$$= \sum_{t=1}^{r} \log p(q^t | q^{t-1}; a, \gamma) + \sum_{t=1}^{r} \log p(\mathcal{S}^t | q^t; \eta) + C$$

Algorithm 2 EM Algorithm for Parameters Learning

Input: $\mathbf{S^r}$;
Output: $\boldsymbol{\theta}$;
1: Initialize $\boldsymbol{\theta^0}$;
2: $k \leftarrow 0$;
3: **while** $\boldsymbol{\theta}$ not converge **do**
4: Compute $Q(\boldsymbol{\theta}, \boldsymbol{\theta^k}) \triangleq \mathbb{E}_{\mathbf{Q^r} \sim p(\mathbf{Q^r} | \mathbf{S^r}; \boldsymbol{\theta^k})}[L(\mathbf{S^r}, \mathbf{Q^r}; \boldsymbol{\theta})]$;
5: $\boldsymbol{\theta^{k+1}} \leftarrow \arg\max_{\boldsymbol{\theta}} Q(\boldsymbol{\theta}, \boldsymbol{\theta^k})$;
6: $k \leftarrow k + 1$;
7: **end while**
8: **return** $\boldsymbol{\theta^k}$;

- Quality inference:
  Linear Dynamical Systems (LDS)
  - Updating formula:

$$\widehat{\mu}^r = \frac{a\eta}{NK + \eta} \widehat{\mu}^{r-1} + \frac{K}{NK + \eta} S, \quad (17)$$

$$\widehat{\sigma}^r = \frac{K\eta}{NK + \eta}, \quad (18)$$

where $K = a^2 \widehat{\sigma}^{r-1} + \gamma$ *(here* a^2 *means* a *squared),* $N = |\mathcal{S}^r|$ *and* $S = \sum_{s_j^r \in \mathcal{S}^r} s_j^r$. *The mean of* $\alpha(q_i^{r+1})$ *(i.e., the estimated quality for the worker in run* $r + 1$*) is*

$$\mu^{r+1} = a\widehat{\mu}^r. \quad (19)$$

Algorithm 3 MELODY Design for Quality Updating

Input: $\widehat{\mu}^{r-1}, \widehat{\sigma}^{r-1}, \mathbf{S^r}$;
Output: $\widehat{\mu}^r, \widehat{\sigma}^r, \mu^{r+1}$;
1: **if** i is a new comer **then**
2: $\widehat{\mu}^r \leftarrow \widehat{\mu}^0, \widehat{\sigma}^r \leftarrow \widehat{\sigma}^0, \mu^{r+1} \leftarrow a\widehat{\mu}^0$;
3: **else**
4: Update $\widehat{\mu}^r$, $\widehat{\sigma}^r$, and μ^{r+1} according to Eq. (17), (18), and (19) respectively;
5: **end if**
6: **if** $\boldsymbol{\theta}$ not updated for T runs **then**
7: $\boldsymbol{\theta} \leftarrow$ Algorithm2($\mathbf{S^r}$);
8: **end if**
9: **return** $\widehat{\mu}^r, \widehat{\sigma}^r, \mu^{r+1}$;

MELODY Framework

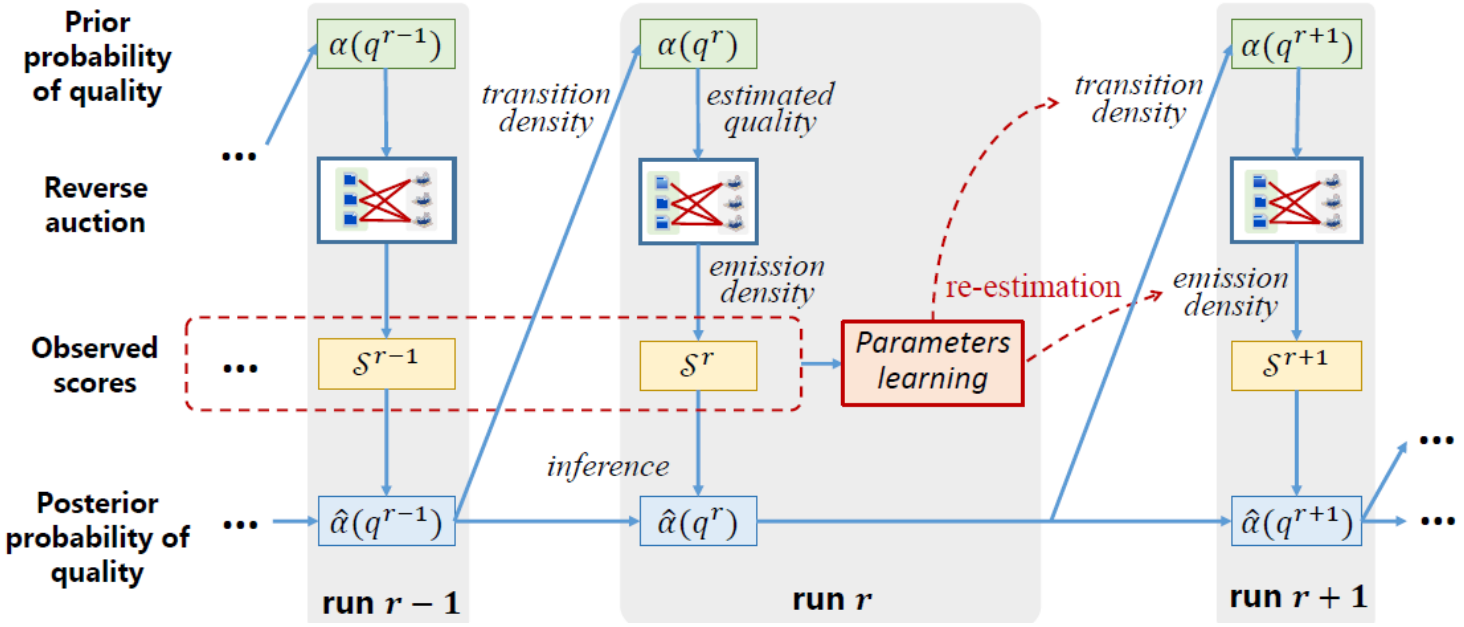

Performance Evaluation (1/3)

- Competitiveness

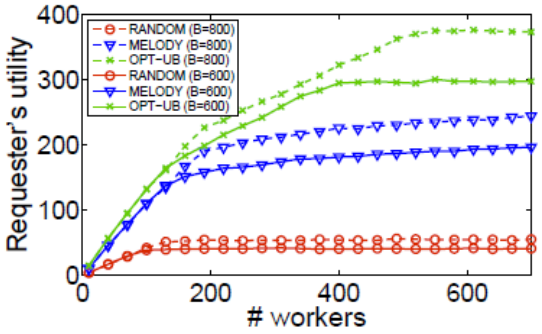

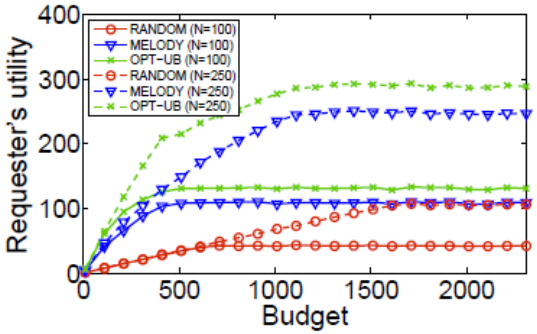

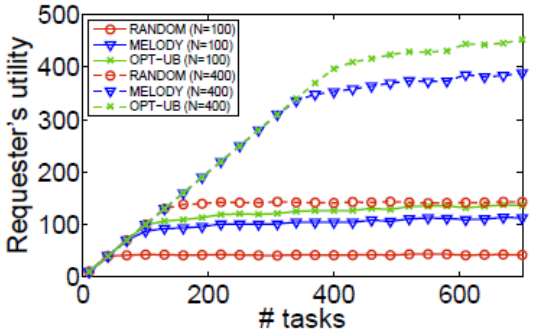

- Individual rationality check

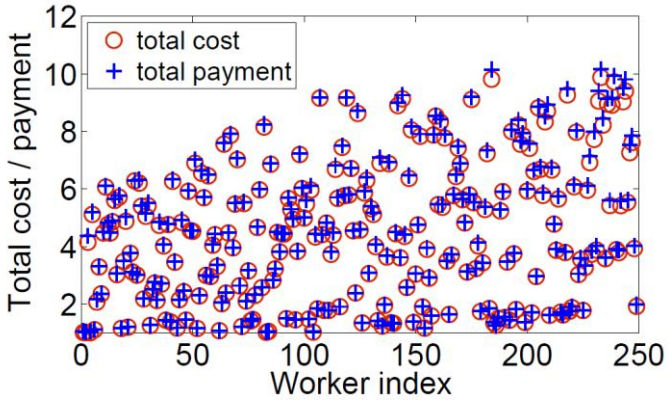

- Budget feasibility check

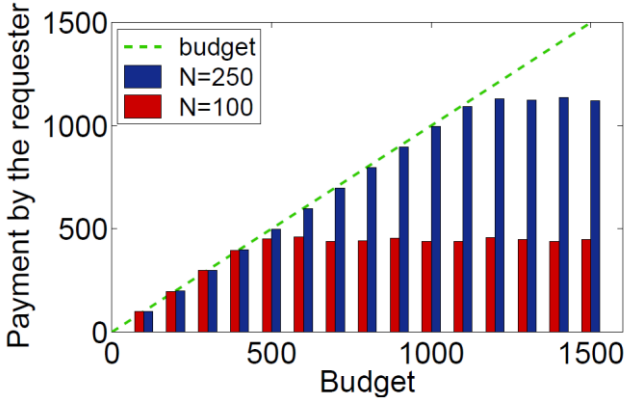

Performance Evaluation (2/3)

- Truthfulness check

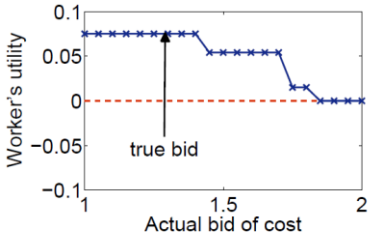

(a) Cost-truthfulness of a winner i

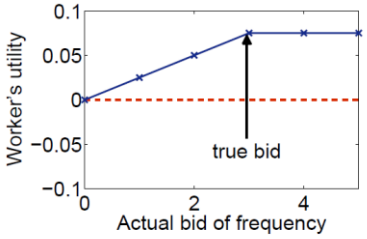

(b) Frequency-truthfulness of a winner i

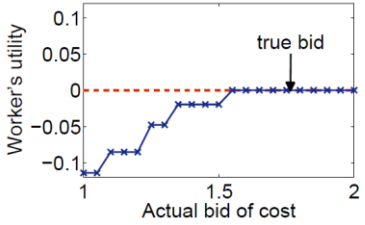

(c) Cost-truthfulness of a loser j

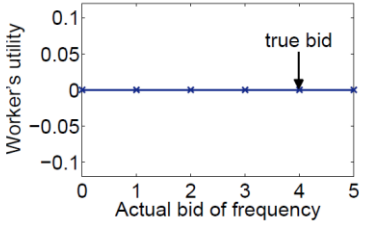

(d) Frequency-truthfulness of a loser j

- Computational efficiency check

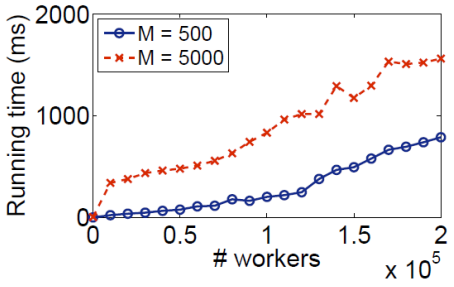

(a) Running time changing with the number of workers

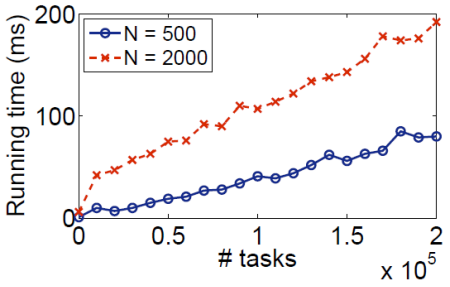

(b) Running time changing with the number of tasks

Performance Evaluation (3/3)

- Long-term quality awareness

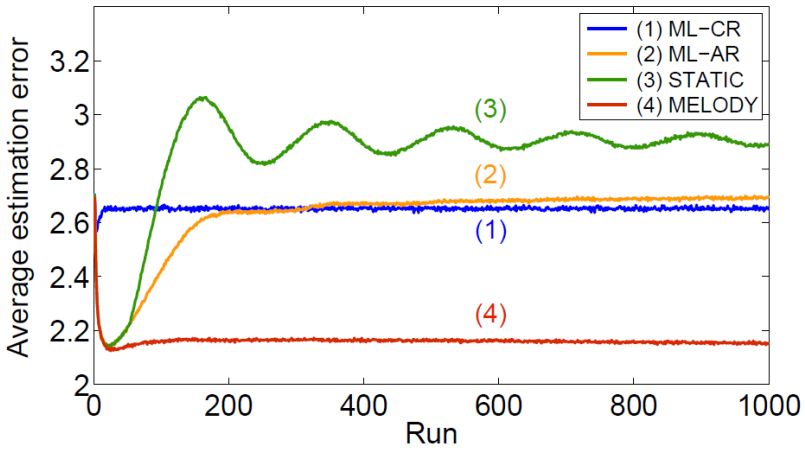

(a) Average estimation error of quality per run

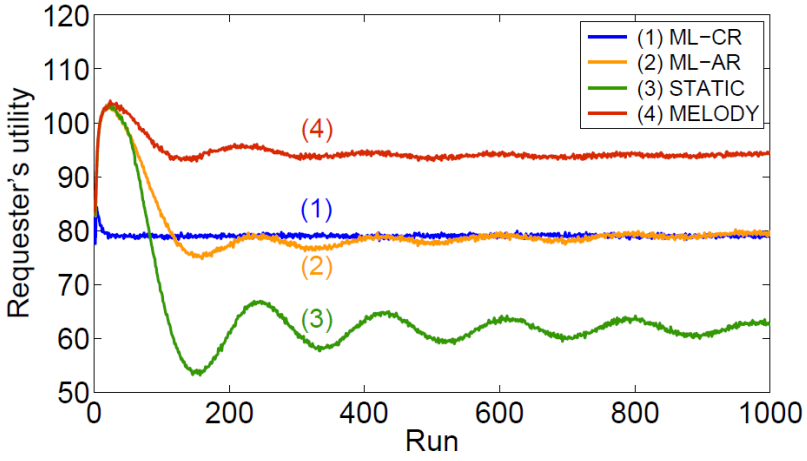

(b) Requester's utility per run

Q & A

Thanks!