

GraphGAN: Graph Representation Learning with Generative Adversarial Nets

Hongwei Wang^{1,2}, Jia Wang³, Jialin Wang⁴, Miao Zhao³,
Weinan Zhang¹, Fuzheng Zhang², Xing Xie², Minyi Guo¹

¹ Shanghai Jiao Tong University, ² Microsoft Research Asia,

³ The Hong Kong Polytechnic University,

⁴ Huazhong University of Science and Technology

January 10, 2018

About Me

❑ Hongwei Wang (王鸿伟)

- ❑ Ph.D. candidate at Shanghai Jiao Tong University
- ❑ Full-time intern at Microsoft Research Asia

❑ Advisor

- ❑ Prof. Minyi Guo (SJTU)
- ❑ Dr. Xing Xie and Dr. Baining Guo (MSRA)

❑ Research interests

- ❑ Recommender systems [TCSS 2017] [RecSys 2017] [CIKM 2017] [WWW 2018]
- ❑ Graph representation learning [WSDM 2018] [AAAI 2018a]
- ❑ Machine learning applications [ICDCS 2017] [TPDS 2018] [AAAI 2018b]

❑ Homepage: <https://hwwang55.github.io>



Outline

❑ Introduction to graph representation learning

- ❑ Definition and application
- ❑ Taxonomy
- ❑ Representative work

❑ GraphGAN [AAAI 2018]

❑ GRL applications

- ❑ Recommender systems [WWW 2018]
- ❑ Sentiment prediction [WSDM 2018]

1. A Comprehensive Survey of Graph Embedding: Problems, Techniques and Applications, TKDE 2017
2. <https://github.com/thunlp/NRLPapers>
3. <https://hwwang55.github.io/files/2017-literature-review-on-rs-and-nrl.pdf> (in Chinese)

Outline

☐ Introduction to graph representation learning

- ☐ Definition and application
- ☐ Taxonomy
- ☐ Representative work


☐ GraphGAN [AAAI 2018]

☐ GRL application

- ☐ Recommender systems [WWW 2018]
- ☐ Sentiment prediction [WSDM 2018]

Definition

- ❑ **Graph representation learning** tries to embed each node of a graph into a low-dimensional vector space, which preserves the structural similarities or distances among the nodes in the original graph
- ❑ Also known as **network embedding** / **graph embedding** / **network representation learning**



Application

❑ **Graph representation learning** can benefit a wide range of real-world applications:

- ❑ Link prediction (Gao, Denoyer, and Gallinari, CIKM 2011)
- ❑ Node classification (Tang, Aggarwal, and Liu, SDM 2016)
- ❑ Recommendation (Yu et al., WSDM 2014)
- ❑ Visualization (Maaten and Hinton, JMLR 2008)
- ❑ Knowledge graph representation (Lin et al., AAAI 2015)
- ❑ Clustering (Tian et al., AAAI 2014)
- ❑ Text embedding (Tang, Qu, and Mei, KDD 2015)
- ❑ Social network analysis (Liu et al., IJCAI 2016)

Taxonomy (1/3)


Input

- ☐ Homogeneous graph (e.g., citation network)
 - ☐ Weighted / Unweighted
 - ☐ Directed / Undirected
 - ☐ Signed / Unsigned
- ☐ Heterogeneous graph
 - ☐ Multimedia network
 - ☐ Knowledge graph
- ☐ Graph with side information
 - ☐ Node/edge label (categorical)
 - ☐ Node/edge attribute (discrete or continuous)
 - ☐ Node feature (e.g., texts)
- ☐ Graph transformed from non-relational data
 - ☐ Manifold learning

Taxonomy (2/3)

Output

- ❑ Node embedding (the most common case)
- ❑ Edge embedding
 - ❑ Relations in knowledge graph
 - ❑ Link prediction
- ❑ Sub-graph embedding
 - ❑ Substructure embedding
 - ❑ Community embedding
- ❑ Whole-graph embedding
 - ❑ Multiple small graphs, e.g., molecule, protein

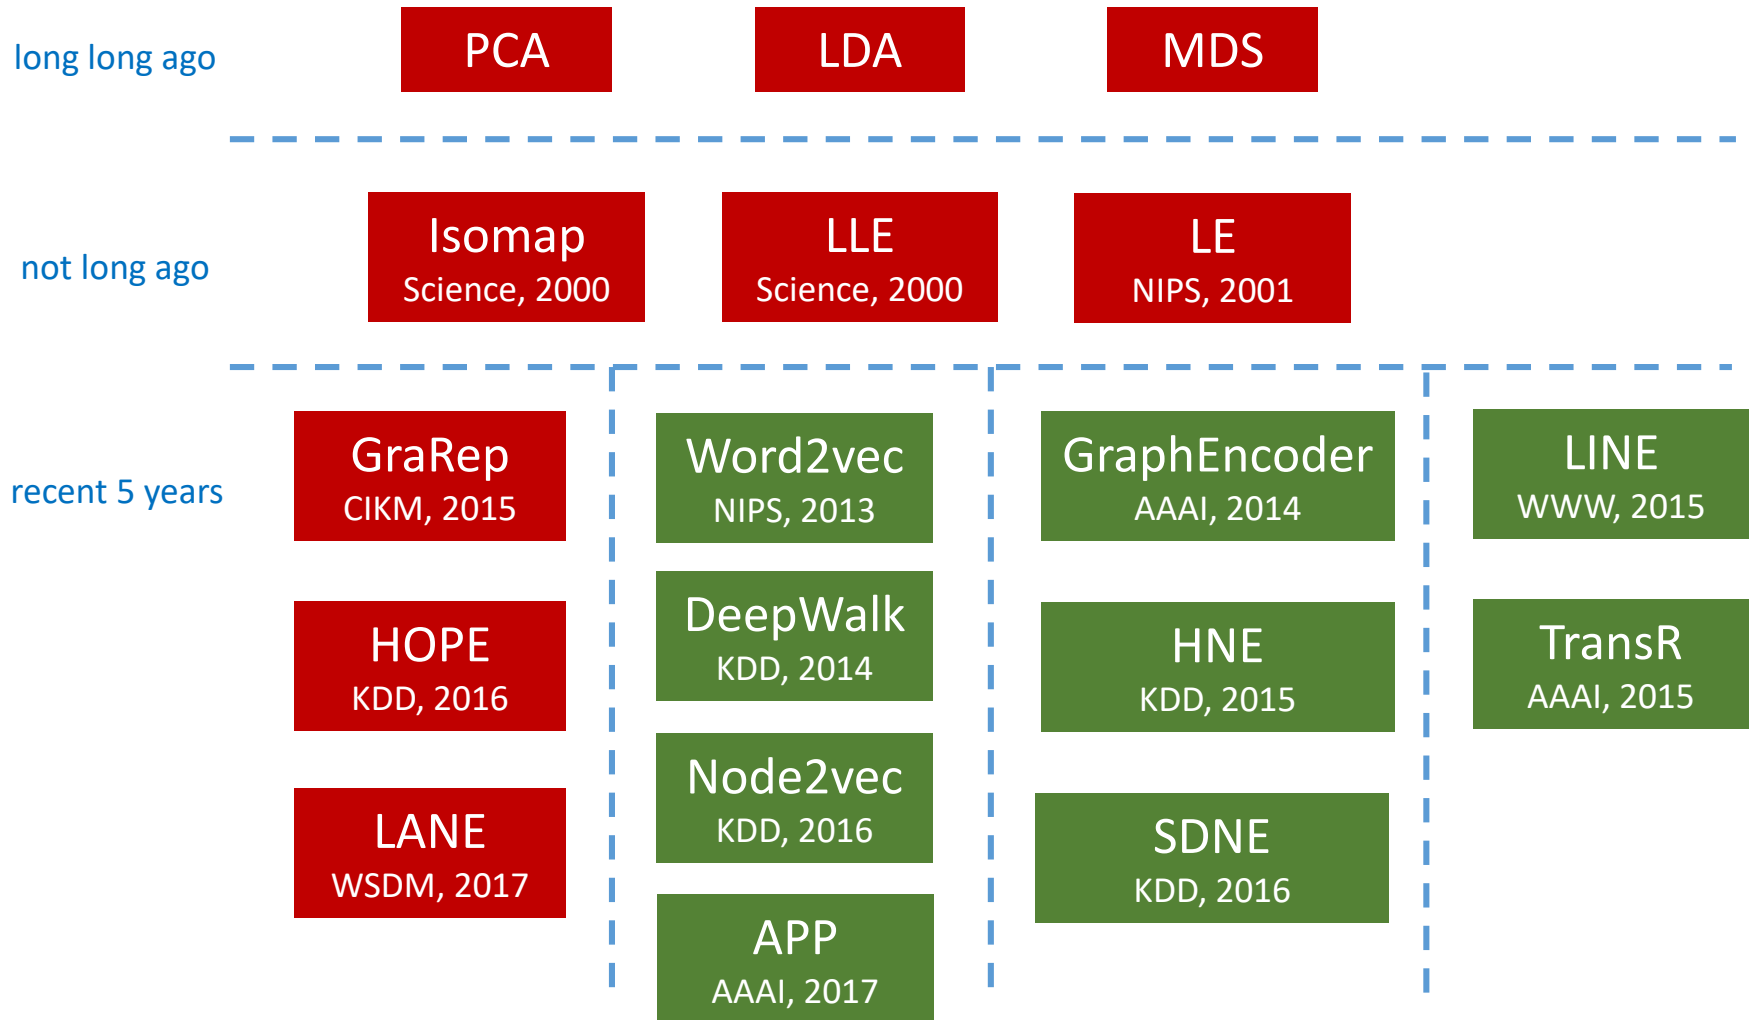


Taxonomy (3/3)

Method

- ☐ Matrix factorization
 - ☐ Singular value decomposition
 - ☐ Spectral decomposition (eigen-decomposition)
- ☐ Random walk
- ☐ Deep learning
 - ☐ Auto-encoder
 - ☐ Convolutional neural network
- ☐ Self-defined loss
 - ☐ Maximizing edge reconstruction probability
 - ☐ Minimizing distance-based loss
 - ☐ Minimizing margin-based ranking loss


Representative Work



Linear Discriminant Analysis

- ❑ Suppose binary classification
- ❑ $D = \{(x_i, y_i)\}$, μ_i : mean of data of the i -th class, Σ_i : covariance matrix of data of the i -th class
- ❑ Make projected covariance matrix as small as possible, while make projected distance between the mean of two classes as large as possible
- ❑ maximize

$$\begin{aligned} J &= \frac{\|w^T \mu_0 - w^T \mu_1\|_2^2}{w^T \Sigma_0 w + w^T \Sigma_1 w} \\ &= \frac{w^T (\mu_0 - \mu_1) (\mu_0 - \mu_1)^T w}{w^T (\Sigma_0 + \Sigma_1) w} \end{aligned}$$

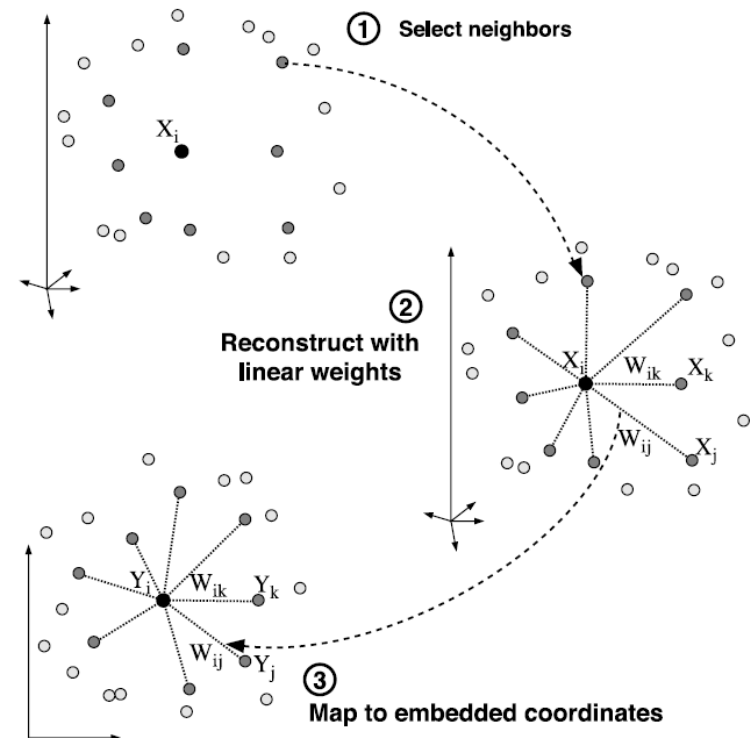



《机器学习》，周志华

Locally Linear Embedding

Introduction

- ❑ An unsupervised learning algorithm that computes low-dimensional, **neighborhood-preserving** embeddings of high-dimensional inputs
- ❑ LLE keeps linear dependency between local instances



Word2vec

❑ Skip-Gram Model

- ❑ Find word representations that are useful for **predicting the surrounding words** in a sentence
- ❑ Maximizing

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

where

$$p(w_O | w_I) = \frac{\exp(v'_{w_O}{}^\top v_{w_I})}{\sum_{w=1}^W \exp(v'_w{}^\top v_{w_I})}$$

- ❑ Negative sampling:

$$\log \sigma(v'_{w_O}{}^\top v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} \left[\log \sigma(-v'_{w_i}{}^\top v_{w_I}) \right]$$

DeepWalk / Node2vec

❑ DeepWalk: Random walk + Word2vec


- ❑ Sample the next node to be visited **uniformly** from the neighbors of current node

$$\underset{\Phi}{\text{minimize}} \quad -\log \Pr \left(\{v_{i-w}, \dots, v_{i+w}\} \setminus v_i \mid \Phi(v_i) \right)$$

- ❑ Hierarchical softmax

❑ Node2vec: Biased random walk + Word2vec

- ❑ Sample the next node to be visited **with bias** from the neighbors of current node




First-order proximity

$$O_1 = d(\hat{p}_1(\cdot, \cdot), p_1(\cdot, \cdot))$$

$$p_1(v_i, v_j) = \frac{1}{1 + \exp(-\vec{u}_i^T \cdot \vec{u}_j)}$$

$$\hat{p}_1(i, j) = \frac{w_{ij}}{W}$$

$$O_1 = - \sum_{(i,j) \in E} w_{ij} \log p_1(v_i, v_j)$$



Second-order proximity

$$O_2 = \sum_{i \in V} \lambda_i d(\hat{p}_2(\cdot | v_i), p_2(\cdot | v_i))$$

$$p_2(v_j | v_i) = \frac{\exp(\vec{u}_j'^T \cdot \vec{u}_i)}{\sum_{k=1}^{|V|} \exp(\vec{u}_k'^T \cdot \vec{u}_i)}$$

$$\hat{p}_2(v_j | v_i) = \frac{w_{ij}}{d_i}$$

$$O_2 = - \sum_{(i,j) \in E} w_{ij} \log p_2(v_j | v_i)$$

TransX

- ❑ Embed **knowledge graph** into a continuous vector space while preserving structural information
- ❑ TransE (NIPS 13):
 - ❑ Ensures $\mathbf{h} + \mathbf{r} \approx \mathbf{t}$ when (h, r, t) holds
- ❑ TransH (AAAI 14):
 - ❑ Ensures $\mathbf{h}_{\perp} + \mathbf{r} \approx \mathbf{t}_{\perp}$ when (h, r, t) holds, where $\mathbf{h}_{\perp} = \mathbf{h} - \mathbf{w}_r^T \mathbf{h} \mathbf{w}_r$ and $\mathbf{t}_{\perp} = \mathbf{t} - \mathbf{w}_r^T \mathbf{t} \mathbf{w}_r$
- ❑ TransR (AAAI 15):
 - ❑ Score function: $f_r(h, t) = \|\mathbf{h} \mathbf{M}_r + \mathbf{r} - \mathbf{t} \mathbf{M}_r\|_2^2$, where \mathbf{M}_r is the projection matrix for relation r .

SDNE

Reconstruction loss term

$$\begin{aligned}\mathcal{L}_{2nd} &= \sum_{i=1}^n \|(\hat{\mathbf{x}}_i - \mathbf{x}_i) \odot \mathbf{b}_i\|_2^2 \\ &= \|(\hat{X} - X) \odot B\|_F^2\end{aligned}$$

Proximity loss term

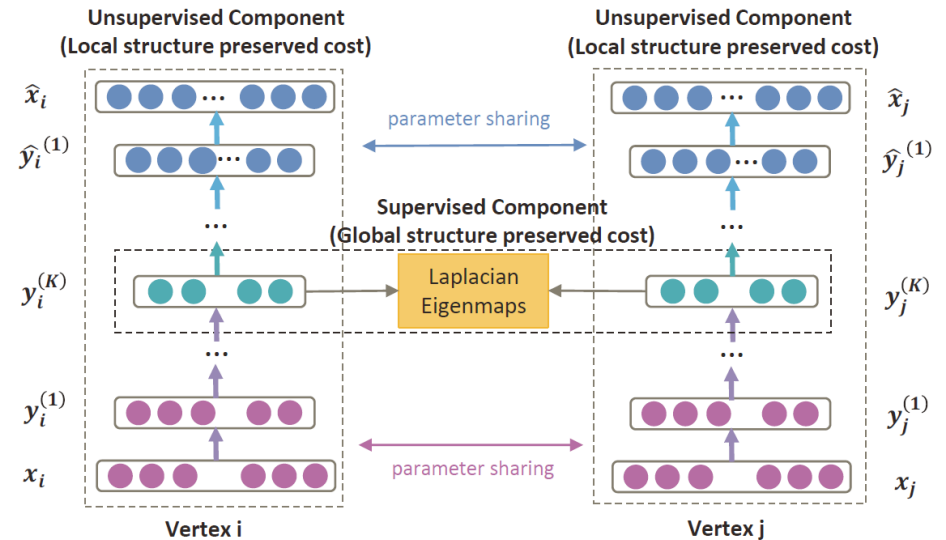
$$\begin{aligned}\mathcal{L}_{1st} &= \sum_{i,j=1}^n s_{i,j} \|y_i^{(K)} - y_j^{(K)}\|_2^2 \\ &= \sum_{i,j=1}^n s_{i,j} \|y_i - y_j\|_2^2\end{aligned}$$

Regularization term

$$\mathcal{L}_{reg} = \frac{1}{2} \sum_{k=1}^K (\|W^{(k)}\|_F^2 + \|\hat{W}^{(k)}\|_F^2)$$

Loss function

$$\mathcal{L}_{mix} = \mathcal{L}_{2nd} + \alpha \mathcal{L}_{1st} + \nu \mathcal{L}_{reg}$$



Outline

☐ Introduction to graph representation learning

- ☐ Definition and application
- ☐ Taxonomy
- ☐ Representative work

☐ **GraphGAN** [AAAI 2018]


☐ GRL application

- ☐ Recommender systems [WWW 2018]
- ☐ Sentiment prediction [WSDM 2018]


Motivation (1/3)

Generative Model

- ❑ Generative graph representation learning model assumes an **underlying true connectivity distribution** $p_{true}(v|v_c)$ for each vertex v_c
 - ❑ The edges can be viewed as observed samples generated by $p_{true}(v|v_c)$
 - ❑ Vertex embeddings are learned by maximizing the likelihood of edges
 - ❑ E.g., DeepWalk (KDD 2014) and node2vec (KDD 2016)



Original graph




$p_{true}(v|v_c)$

Motivation (2/3)

Discriminative Model

- ❑ Discriminative graph representation learning model aim to learn a **classifier** for predicting the existence of edges directly
 - ❑ Consider two vertices v_i and v_j jointly as features
 - ❑ Predict the probability of an edge existing between them, i.e., $p(\text{edge}|v_i, v_j)$
 - ❑ E.g., SDNE (KDD 2016) and PPNE (DASFAA, 2017)



$$p(\text{edge}|v_i, v_j) = 0.8$$

$$p(\text{edge}|v_i, v_k) = 0.3$$

.....

Motivation (3/3)

G + D ?

- ❑ Generative and discriminative models are two sides of the same coin
- ❑ LINE (WWW 2015) has tried to combine these two objectives
- ❑ Generative adversarial nets (GAN) have received a great deal of attention
 - ❑ GAN designs a game-theoretical minimax game to combine G and D
 - ❑ GAN achieves success in various applications:
 - ❑ image generation (Denton et al., NIPS 2015)
 - ❑ sequence generation (Yu et al., AACL 2017)
 - ❑ dialogue generation (Li et al., arXiv 2017)
 - ❑ information retrieval (Wang et al., SIGIR 2017)
 - ❑ domain adaption (Zhang, Barzilay, and Jaakkola, arXiv 2017)
- ❑ We propose **GraphGAN**, a framework that unifies generative and discriminative thinking for graph representation learning

The Minimax Game

- ❑ $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, $\mathcal{V} = \{v_1, \dots, v_V\}$, $\mathcal{E} = \{e_{ij}\}_{i,j=1}^V$
- ❑ $\mathcal{N}(v_c)$: set of neighbors of v_c
- ❑ $p_{true}(v_c)$: underlying true connectivity distribution for v_c

- ❑ The objective of GraphGAN is to learn the following two models:
 - ❑ $G(v|v_c; \theta_G)$ which tries to approximate $p_{true}(v_c)$
 - ❑ $D(v, v_c; \theta_D)$ which aims to discriminate the connectivity for the vertex pair (v, v_c)

- ❑ The two-player minimax game:

$$\min_{\theta_G} \max_{\theta_D} V(G, D) = \sum_{c=1}^V \left(\mathbb{E}_{v \sim p_{true}(\cdot|v_c)} [\log D(v, v_c; \theta_D)] + \right. \\ \left. \mathbb{E}_{v \sim G(\cdot|v_c; \theta_G)} [\log (1 - D(v, v_c; \theta_D))] \right) \quad (1)$$

Implementation & Optimization of D

$$\min_{\theta_G} \max_{\theta_D} V(G, D) = \sum_{c=1}^V \left(\mathbb{E}_{v \sim p_{\text{true}}(\cdot | v_c)} [\log D(v, v_c; \theta_D)] + \mathbb{E}_{v \sim G(\cdot | v_c; \theta_G)} [\log (1 - D(v, v_c; \theta_D))] \right) \quad (1)$$

□ Implementation of D:

$$D(v, v_c; \theta_D) = \sigma(\mathbf{d}_v^\top \mathbf{d}_{v_c}) = \frac{1}{1 + \exp(-\mathbf{d}_v^\top \mathbf{d}_{v_c})}, \quad (2)$$

where $\mathbf{d}_v, \mathbf{d}_{v_c} \in \mathbb{R}^k$ are the k-dimensional vectors of v and v_c for D

□ Gradient of $V(G, D)$ w.r.t θ_D :

$$\nabla_{\theta_D} V(G, D) = \begin{cases} \nabla_{\mathbf{d}_v, \mathbf{d}_{v_c}} \log D(v, v_c; \theta_D), & \text{if } v \sim p_{\text{true}}; \\ \nabla_{\mathbf{d}_v, \mathbf{d}_{v_c}} (1 - \log D(v, v_c; \theta_D)), & \text{if } v \sim G. \end{cases} \quad (3)$$

Optimization of G

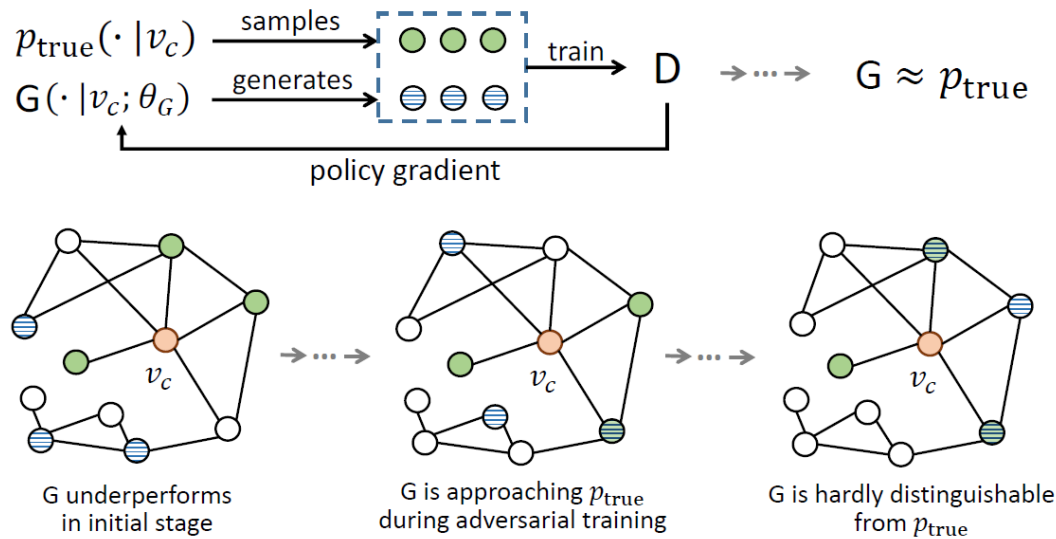
$$\min_{\theta_G} \max_{\theta_D} V(G, D) = \sum_{c=1}^V \left(\mathbb{E}_{v \sim p_{\text{true}}(\cdot|v_c)} [\log D(v, v_c; \theta_D)] + \mathbb{E}_{v \sim G(\cdot|v_c; \theta_G)} [\log (1 - D(v, v_c; \theta_D))] \right) \quad (1)$$

□ Gradient of $V(G, D)$ w.r.t θ_G (policy gradient):

$$\begin{aligned} & \nabla_{\theta_G} V(G, D) \\ &= \nabla_{\theta_G} \sum_{c=1}^V \mathbb{E}_{v \sim G(\cdot|v_c; \theta_G)} [\log (1 - D(v, v_c; \theta_D))] \\ &= \sum_{c=1}^V \sum_{i=1}^N \nabla_{\theta_G} G(v_i|v_c; \theta_G) \log (1 - D(v_i, v_c; \theta_D)) \\ &= \sum_{c=1}^V \sum_{i=1}^N G(v_i|v_c; \theta_G) \nabla_{\theta_G} \log G(v_i|v_c; \theta_G) \log (1 - D(v_i, v_c; \theta_D)) \\ &= \sum_{c=1}^V \mathbb{E}_{v \sim G(\cdot|v_c; \theta_G)} [\nabla_{\theta_G} \log G(v|v_c; \theta_G) \log (1 - D(v, v_c; \theta_D))] \end{aligned} \quad (4)$$

GraphGAN Framework

$$\min_{\theta_G} \max_{\theta_D} V(G, D) = \sum_{c=1}^V \left(\mathbb{E}_{v \sim p_{\text{true}}(\cdot | v_c)} [\log D(v, v_c; \theta_D)] + \mathbb{E}_{v \sim G(\cdot | v_c; \theta_G)} [\log (1 - D(v, v_c; \theta_D))] \right) \quad (1)$$



Implementation of G

❑ Softmax?

- ❑ Computationally inefficient
- ❑ Graph-structure-unaware


$$G(v|v_c; \theta_G) = \frac{\exp(\mathbf{g}_v^\top \mathbf{g}_{v_c})}{\sum_{v \neq v_c} \exp(\mathbf{g}_v^\top \mathbf{g}_{v_c})}$$

where $\mathbf{g}_v, \mathbf{g}_{v_c} \in \mathbb{R}^k$ are the k -dimensional vectors of v and v_c for G

❑ Hierarchical softmax?

- ❑ Graph-structure-unaware

$$p(w|w_I) = \prod_{j=1}^{L(w)-1} \sigma \left(\mathbb{I}[n(w, j+1) = \text{ch}(n(w, j))] \cdot v'_{n(w, j)}{}^\top v_{w_I} \right)$$



❑ Negative sampling?

- ❑ Not a valid probability distribution
- ❑ graph-structure-unaware

$$\log \sigma(v'_{w_O}{}^\top v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} \left[\log \sigma(-v'_{w_i}{}^\top v_{w_I}) \right]$$

Graph Softmax (1/5)

Objectives

- ❑ The design of graph softmax should satisfy the following three properties:
 - ❑ **Normalized:** The generator should produce a valid probability distribution, i.e.,
$$\sum_{v \neq v_c} G(v|v_c; \theta_G) = 1$$
 - ❑ **Graph-structure-aware:** The generator should take advantage of the structural information of a graph
 - ❑ **Computationally efficient:** The computation of $G(v|v_c; \theta_G)$ should only involve a small number of vertices in the graph

Graph Softmax (2/5)

Design

❑ **Breadth First Search (BFS)** on \mathcal{G} from every vertex v_c

❑ **BFS-tree** T_c rooted at v_c

❑ For a given vertex v and one of its neighbors $v_i \in \mathcal{N}_c(v)$, the **relevance probability** of v_i given v as

$$p_c(v_i|v) = \frac{\exp(\mathbf{g}_{v_i}^\top \mathbf{g}_v)}{\sum_{v_j \in \mathcal{N}_c(v)} \exp(\mathbf{g}_{v_j}^\top \mathbf{g}_v)}, \quad (6)$$

where $\mathcal{N}_c(v)$ is the set of neighbors of v in T_c

❑ **Graph softmax**

$$G(v|v_c; \theta_G) \triangleq \left(\prod_{j=1}^m p_c(v_{r_j}|v_{r_{j-1}}) \right) \cdot p_c(v_{r_m}|v_{r_{m-1}}), \quad (7)$$

given the unique path from v_c to v in tree T_c : $P_{v_c \rightarrow v} = (v_{r_0}, v_1, \dots, v_{r_m})$, where $v_{r_0} = v_c$ and $v_{r_m} = v$

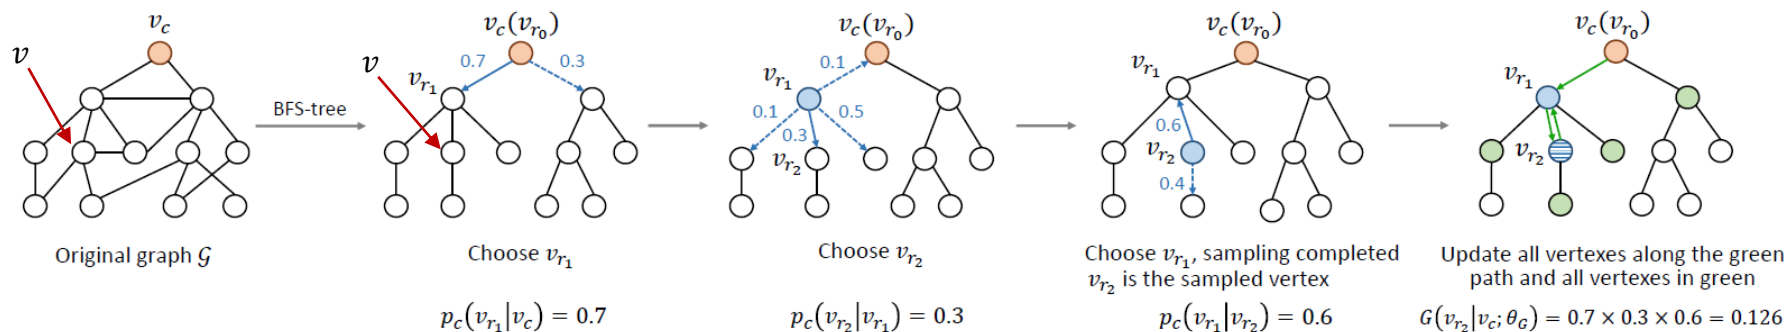
Graph Softmax (3/5)

Design

□ Graph softmax

$$G(v|v_c; \theta_G) \triangleq \left(\prod_{j=1}^m p_c(v_{r_j}|v_{r_{j-1}}) \right) \cdot p_c(v_{r_m-1}|v_{r_m}), \quad (7)$$

given the unique path from v_c to v in tree T_c : $P_{v_c \rightarrow v} = (v_{r_0}, v_1, \dots, v_{r_m})$, where $v_{r_0} = v_c$ and $v_{r_m} = v$



Graph Softmax (4/5)

Properties

- $\sum_{v \neq v_c} G(v|v_c; \theta_G) = 1$ in graph softmax
- In graph softmax, $G(v|v_c; \theta_G)$ decreases exponentially with the increase of the shortest distance between v and v_c in original graph \mathcal{G}
- In graph softmax, calculation of $G(v|v_c; \theta_G)$ depends on $O(d \log V)$ vertices, where d is average degree of vertices and V is the number of vertices in graph \mathcal{G}

Graph Softmax (5/5)

Generating Strategy

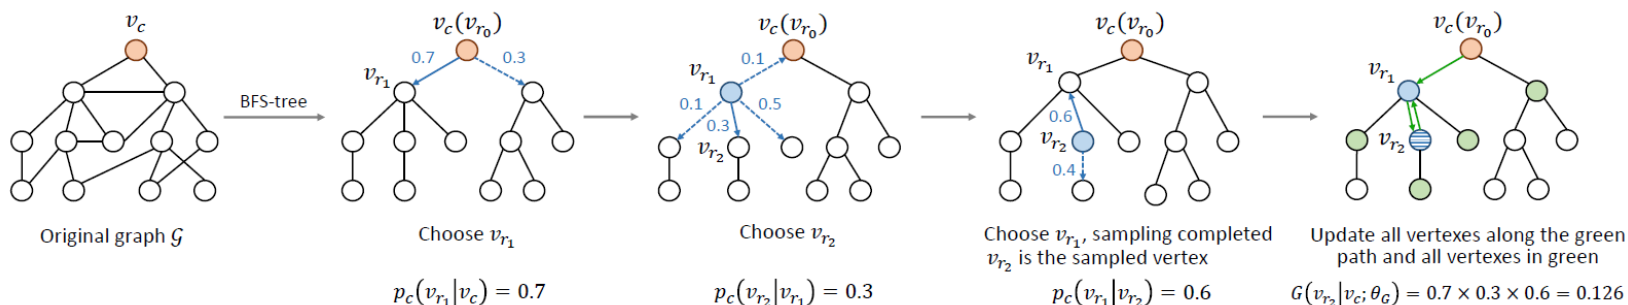
Algorithm 1 Online generating strategy for the generator

Input: BFS-tree T_c , representation vectors $\{\mathbf{g}_i\}_{i \in \mathcal{V}}$

Output: generated sample v_{gen}

```

1:  $v_{pre} \leftarrow v_c, v_{cur} \leftarrow v_c$ ;
2: while true do
3:   Randomly select  $v_i$  proportionally to  $p_c(v_i|v_{cur})$  in Eq.
   (6);
4:   if  $v_i = v_{pre}$  then
5:      $v_{gen} \leftarrow v_{cur}$ ;
6:     return  $v_{gen}$ 
7:   else
8:      $v_{pre} \leftarrow v_{cur}, v_{cur} \leftarrow v_i$ ;
9:   end if
10: end while
  
```



GraphGAN Algorithm

Algorithm 2 GraphGAN framework

Input: dimension of embedding k , size of generating samples s , size of discriminating samples t

Output: generator $G(v|v_c; \theta_G)$, discriminator $D(v, v_c; \theta_D)$

- 1: Initialize and pre-train $G(v|v_c; \theta_G)$ and $D(v, v_c; \theta_D)$;
 - 2: Construct BFS-tree T_c for all $v_c \in \mathcal{V}$;
 - 3: **while** GraphGAN not converge **do**
 - 4: **for** G-steps **do**
 - 5: $G(v|v_c; \theta_G)$ generates s vertices for each vertex v_c according to Algorithm 1;
 - 6: Update θ_G according to Eq. (4), (6) and (7);
 - 7: **end for**
 - 8: **for** D-steps **do**
 - 9: Sample t positive vertices from ground truth and t negative vertices from $G(v|v_c; \theta_G)$ for each vertex v_c ;
 - 10: Update θ_D according to Eq. (2) and (3);
 - 11: **end for**
 - 12: **end while**
 - 13: **return** $G(v|v_c; \theta_G)$ and $D(v, v_c; \theta_D)$
-

Experiments (1/3)

Datasets

- ❑ arXiv-AstroPh: 18,772 vertices and 198,110 edges
- ❑ arXiv-GrQc: 5,242 vertices and 14,496 edges
- ❑ BlogCatalog: 10,312 vertices, 333,982 edges and 39 labels
- ❑ Wikipedia: 4,777 vertices, 184,812 edges and 40 labels
- ❑ MovieLens-1M: 6,040 users and 3,706 movies

Baselines

- ❑ DeepWalk (KDD 2014)
- ❑ LINE (WWW 2015)
- ❑ Node2vec (KDD 2016)
- ❑ Struc2vec (KDD 2017)

Experiments (2/3)

Link Prediction

□ Learning curves




Fig. 4: Learning curves of the generator and the discriminator of GraphGAN on arXiv-GrQc in link prediction.

□ Results

TABLE 1: Accuracy and Macro-F1 on arXiv-AstroPh and arXiv-GrQc in link prediction.

Model	arXiv-AstroPh		arXiv-GrQc	
	Accuracy	Macro-F1	Accuracy	Macro-F1
DeepWalk	0.841	0.839	0.803	0.812
LINE	0.820	0.814	0.764	0.761
Node2vec	0.845	0.854	0.844	0.842
Struc2vec	0.821	0.810	0.780	0.776
GraphGAN	0.855	0.859	0.849	0.853


Experiments (3/3)

Node Classification


TABLE 2: Accuracy and Macro-F1 on BlogCatalog and Wikipedia in node classification.

Model	BlogCatalog		Wikipedia	
	Accuracy	Macro-F1	Accuracy	Macro-F1
DeepWalk	0.225	0.214	0.194	0.183
LINE	0.205	0.192	0.175	0.164
Node2vec	0.215	0.206	0.191	0.179
Struc2vec	0.228	0.216	0.211	0.190
GraphGAN	0.232	0.221	0.213	0.194

Recommendation



(a) Precision@K



(b) Recall@K

Fig. 5: Precision@K and Recall@K on MovieLens-1M in recommendation.

Summary

- ❑ We propose **GraphGAN**, a novel framework that unifies generative and discriminative thinking for graph representation learning
 - ❑ Generator $G(v|v_c)$ tries to fit $p_{true}(v|v_c)$ as much as possible
 - ❑ Discriminator $D(v, v_c)$ tries to tell whether an edge exists between v and v_c

- ❑ G and D act as two players in a **minimax game**:
 - ❑ G tries to produce the most indistinguishable “fake” vertices under guidance provided by D
 - ❑ D tries to draw a clear line between the ground truth and “counterfeits” to avoid being fooled by G

- ❑ We propose **graph softmax** as the implementation of G
 - ❑ Graph softmax overcomes the limitations of softmax and hierarchical softmax
 - ❑ Graph softmax satisfies the properties of **normalization**, **graph structure awareness** and **computational efficiency**

Outline

❑ Introduction to graph representation learning

- ❑ Definition and application
- ❑ Taxonomy
- ❑ Representative work

❑ GraphGAN [AAAI 2018]

❑ GRL application

- ❑ Recommender systems [WWW 2018]
- ❑ Sentiment prediction [WSDM 2018]

DKN (WWW 2018)

❑ DKN: Deep Knowledge-Aware Network for News Recommendation

- ❑ Learning knowledge graph representations by TransX
- ❑ A CNN framework for combining word embedding and entity embedding
- ❑ Attention-based CTR prediction

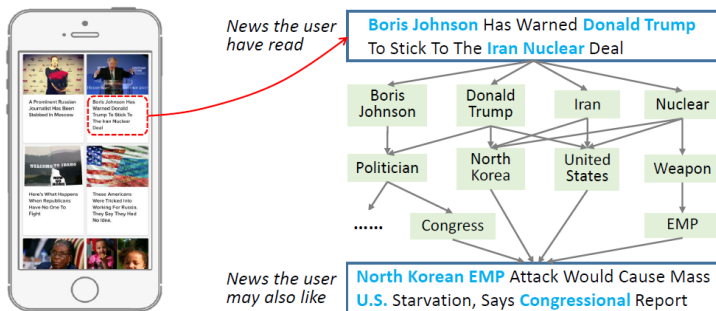


Fig. 1: Knowledge graph in news recommendation

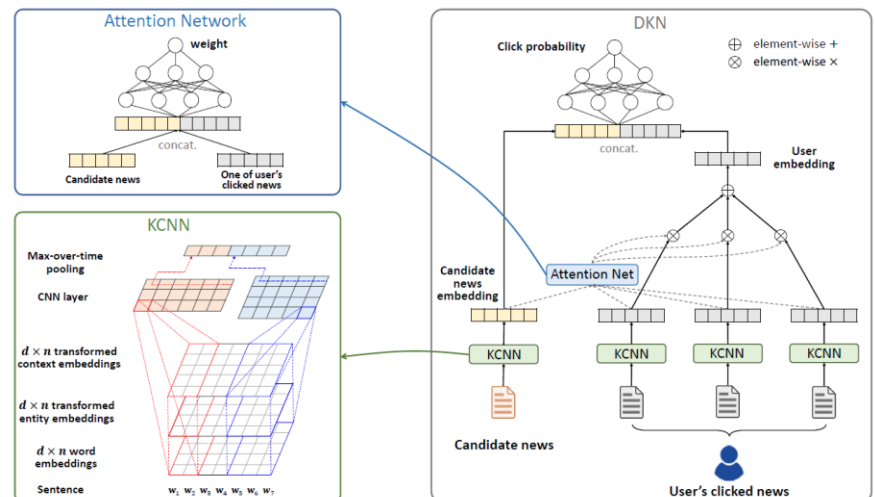


Fig. 2: DKN framework

SHINE (WSDM 2018)

SHINE: Signed Heterogeneous Information Network Embedding for Sentiment Link Prediction

- Sign: positive/negative sentiment link
- Heterogeneity: sentiment network, social network, knowledge graph
- Auto-encoder based framework




Fig. 1: Signed heterogeneous networks in sentiment prediction




Fig. 2: Auto-encoder for sentiment network representation learning

SHINE (WSDM 2018)

SHINE: Signed Heterogeneous Information Network Embedding for Sentiment Link Prediction

- Sign: positive/negative sentiment link
- Heterogeneity: sentiment network, social network, knowledge graph
- Auto-encoder based framework

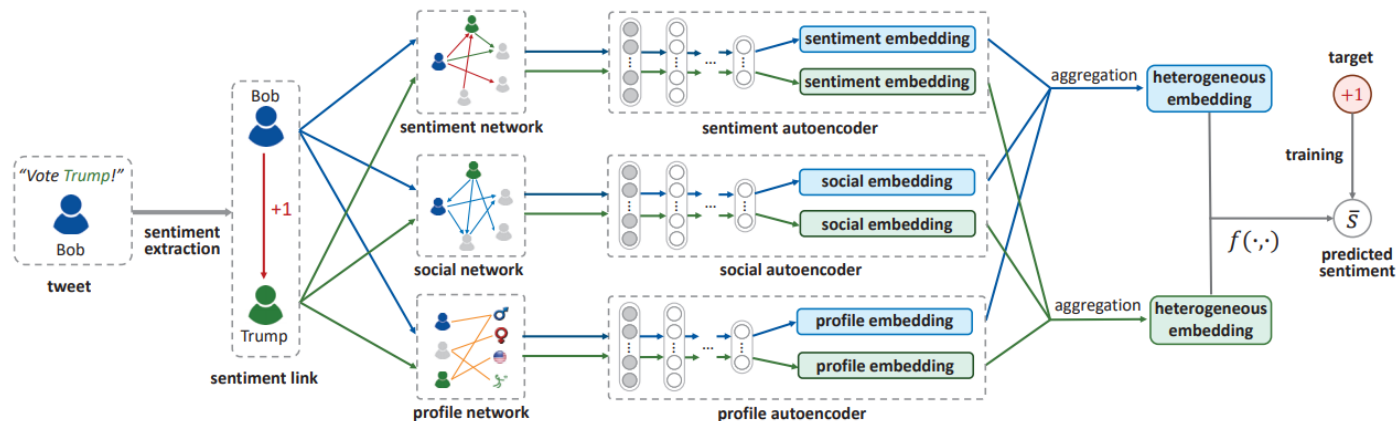


Fig. 3: SHINE framework

Q & A

Thanks!

Visit <https://hwwang55.github.io> for full papers, slides, code and more information